

Performance evaluation of Terapixel rendering in Cloud (Super)computing

Submitive Assignment
Rishabh Malik(220512619)

1. Introduction

Terapixel images, which are composed of a vast number of pixels, can present information to stakeholders in an effective manner, but they require a high amount of computational power. To reduce costs and speed up the process, it is essential to use cloud-based supercomputing resources to process the data collected during the creation of a terapixel image of Newcastle Airport. The issue at hand is how to access the necessary supercomputing resources for the calculation of this large-scale image in the cloud. The objective is to analyze the data from the application checkpoint and system metrics output to improve the cost-efficiency of the process.

2. Business Understanding

The business understanding phase focuses on understanding the requirement from the business perspective of the project. This is an important aspect of business for this project as it helps organizations understand its capabilities and limitations of their cloud-based supercomputing infrastructure in handling large-scale image rendering tasks.

This information can be used to make informed decisions about the use of cloud computing resources for large-scale image rendering, such as whether to invest in additional resources or optimize existing ones. Additionally, it can also help organizations identify bottlenecks in their workflow and make adjustments to improve the performance and cost-efficiency of the process.

Furthermore, the performance evaluation can also be used to compare the performance of different cloud providers or different configurations, which can be useful for organizations to choose the best solution for their specific needs. In short, this helps organizations to optimize their cloud-based supercomputing infrastructure to handle large-scale image rendering tasks efficiently, make informed decisions about the use of cloud resources, and improve the overall performance and cost-efficiency of the process.

Through this analysis we will gain insights into the following:

- Which event types dominate task runtimes?
- What is the interplay between GPU temperature and performance?
- What is the interplay between increased power draw and render time?
- Can we identify particular GPU cards (based on their serial numbers) whose performance differs from other cards? (i.e. perpetually slow cards).

3. Data Understanding

Adding to the foundation of Business Understanding, it drives the focus to identify, collect, and analyze the data sets that can help you accomplish the project goals. Here we acquire the necessary data and load it into our analysis tool and get a field overview of the dataset and its format, number of records, and other fields.

The data presented in this report was obtained from an experiment that employed a massive 1024 GPU setup. The rendering process was broken down into three distinct levels, specifically, levels 4, 8, and 12 of the final output. The information provided in the dataset gives an in-depth insight into the timing of the rendering performance, the performance of the GPU, and the exact area of the image that was being handled during each job.

Our data consist of three files:

- application-checkpoint.csv
- gpu.csv
- task-x-y.csv

- application-checkpoint.csv:** This file contains information on the various events that occurred during the rendering process, including the type of event, the timestamp it occurred, and the hostname, taskid, and jobid associated with it.
 - timestamp:** this column records the precise moment the event began or ended.
 - hostname:** this column lists the 1024 unique GPUs used in the operation.
 - eventName:** this column includes 5 different types of events: Tiling, Render, Saving Config, Uploading, and TotalRender.
 - TotalRender:** this represents the overall task.
 - Render:** this is when specific tiles of the image are being rendered.
 - Tiling:** is the post-processing of the rendered tiles.
 - Saving Config:** is when the specific configurations for each task are saved
 - Uploading:** is when the post-processed tiles are uploaded to Azure Blob Storage.
 - eventType:** This column indicates whether the event is starting or ending
- jobid:** this column indicates the specific level of the visualization output (level 12, level 8, and level 4)
- taskid:** this column provides a unique identification for each set of events, including their start and stop timestamps. There will be n*10 unique taskids in the file, where n is the number of rows in the application-checkpoint.csv.
- gpu.csv:** This file contains information about the status of the graphics processing units (GPUs) at a specific point in time.
 - gpuSerial:** this file includes details such as the unique serial number for each of the 1024 GPUs.
 - gpuUID:** This is the system-assigned unique identifier for each GPU.
 - powerDrawWatt:** the amount of power being used by the GPU in watts.
 - gpuTempC:** the temperature of the GPU in degrees Celsius
 - gpuMemUtilPerc:** the percentage of the GPU's cores that are being used.
 - gpuMemUtilPerc:** the percentage of the GPU's memory that is being utilized.
- task-x-y.csv:** This file contains information about the specific sections of an image that are being rendered.
 - x** and **y:** it includes the x and y coordinates of the tiles being processed.
 - level:** this is the level of zoom being used. The visualization is designed to be similar to a "google maps style" map, with 12 levels of zoom available. The data only includes levels 4, 8, and 12, as the intermediate levels are created during the post-processing, or tiling, phase.

4. Data Processing

This phase is also called data munging. Here we will now merge the 3 data files into one to make a meaningful dataset.

- Firstly left join the 2 data files task-x-y.csv and application-checkpoint.csv based on two columns jobid and taskid.
- Secondly inner join to gpu.csv file with the data from the previous step based on column timestamp.

Now on the final data file data cleaning is to be performed, as there are some duplicate values and we can see that the timestamp column is not in the proper format. Hence we delete the duplicates in the final data file. Also, we use the datetime library to format the timestamp column to make it standard time for further calculations.

As we have already understood in the Business Understanding phase we need to calculate the duration of the events, hence we need to perform munge on the timestamp column. Now we create 2 data frames df_start and df_stop. For df_start we take the eventType = "START" and put the timestamp in the as_start_time and similarly for eventType = "STOP" we put the timestamp in stop_time. Merged start and stop times of respective events on taskid and then subtract start and stop time to compute the duration and store it in a duration column. Now the final merged dataset has the duration time as well for further computations.

5. Exploratory Data Analysis

```
In [78]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import datetime
from matplotlib import pyplot
from pandas import Series
import statistics
from statistics import mean
from dateutil import parser
```

```
In [80]: #reading the csv files
gpudf = pd.read_csv('gpu.csv')
appdf = pd.read_csv('application_checkpoints.csv')
taskdf = pd.read_csv('task-x-y.csv')
```

```
In [81]: app_task_df = pd.merge(taskdf, appdf, on = ['jobid', 'taskid'], how='left') #applying left join
app_task_gpu_df = pd.merge(app_task_df, gpudf, on = ['timestamp']) #applying inner join
```

```
In [82]: app_task_gpu_df.head()
```

	taskid	jobid	x	y	level	timestamp	hostname	x_eventName	eventType	hostname_y	gpuSerial	gpuUID	powerC
0	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	116	178	12	2018-11-08T08:07:10.688Z	074591444K046078517041d702622b700001	Tiling	STOP	265232c596814768eaef65a7becf690000Q	32361702812	GPU-82614902-9605-3605-952d364289	GPU-82614902-9605-3605-952d364289
1	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	116	178	12	2018-11-08T08:07:10.688Z	074591444K046078517041d702622b700001	Tiling	STOP	0056a730076643b58977f00d0d882c1a0000Q	325117171574	GPU-48410234-6202-4942-9605-952d364289	GPU-48410234-6202-4942-9605-952d364289
2	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	142	190	12	2018-11-08T08:14:48.852Z	83ea61ac1e5427a3bf7b0d041ecaa700001A	Uploading	START	d8241877c0994572b46e4681e5d144d850000W	323617021323	GPU-8711640-9605-3605-952d364289	GPU-8711640-9605-3605-952d364289
3	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	142	190	12	2018-11-08T08:14:48.852Z	83ea61ac1e5427a3bf7b0d041ecaa700001A	Render	STOP	d8241877c0994572b46e4681e5d144d850000W	323617021323	GPU-8711640-9605-3605-952d364289	GPU-8711640-9605-3605-952d364289
4	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	142	190	12	2018-11-08T08:14:48.852Z	83ea61ac1e5427a3bf7b0d041ecaa700001A	TotalRender	STOP	8b64beeb87b4c3b20539e81075191800000U	323617021463	GPU-8920498-9605-3605-952d364289	GPU-8920498-9605-3605-952d364289

```
In [83]: app_task_gpu_df = app_task_gpu_df.drop_duplicates(subset=None, keep='first', inplace=False) # deleting duplicate rows
app_task_gpu_df.shape
```

```
In [85]: app_task_gpu_df['eventName'] = app_task_gpu_df['timestamp'].apply(lambda x: parser.isoparse(x)) # extracting only time part and storing it in "time" column
app_task_gpu_df.head()
```

```
In [86]: app_task_gpu_df.head()
```

	taskid	jobid	x	y	level	timestamp	hostname	x_eventName	eventType	hostname_y	gpuSerial	gpuUID	powerC
0	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	116	178	12	2018-11-08T08:07:10.688Z	074591444K046078517041d702622b700001	Tiling	STOP	265232c596814768eaef65a7becf690000Q	32361702812	GPU-82614902-9605-3605-952d364289	GPU-82614902-9605-3605-952d364289
1	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	116	178	12	2018-11-08T08:07:10.688Z	074591444K046078517041d702622b700001	Tiling	STOP	0056a730076643b58977f00d0d882c1a0000Q	325117171574	GPU-48410234-6202-4942-9605-952d364289	GPU-48410234-6202-4942-9605-952d364289
2	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	142	190	12	2018-11-08T08:14:48.852Z	83ea61ac1e5427a3bf7b0d041ecaa700001A	Uploading	START	d8241877c0994572b46e4681e5d144d850000W	323617021323	GPU-8711640-9605-3605-952d364289	GPU-8711640-9605-3605-952d364289
3	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	142	190	12	2018-11-08T08:14:48.852Z	83ea61ac1e5427a3bf7b0d041ecaa700001A	Render	STOP	d8241877c0994572b46e4681e5d144d850000W	323617021323	GPU-8711640-9605-3605-952d364289	GPU-8711640-9605-3605-952d364289
4	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	142	190	12	2018-11-08T08:14:48.852Z	83ea61ac1e5427a3bf7b0d041ecaa700001A	TotalRender	STOP	8b64beeb87b4c3b20539e81075191800000U	323617021463	GPU-8920498-9605-3605-952d364289	GPU-8920498-9605-3605-952d364289

```
In [87]: #for calculating the duration (stop time - start time)
df_start = app_task_gpu_df[app_task_gpu_df['eventType'] == "START"]
df_stop = app_task_gpu_df[app_task_gpu_df['eventType'] == "STOP"]
df_start = df_start.rename(columns={"timestamp": "start_time"})
df_stop = df_stop.rename(columns={"timestamp": "stop_time"})
df_start = df_start.drop('eventType', axis = 1)
df_stop = df_stop.drop('eventType', axis = 1)
```

```
# convert columns to datetime
df_start['start_time'] = pd.to_datetime(df_start['start_time'])
df_stop['stop_time'] = pd.to_datetime(df_stop['stop_time'])
df = pd.merge(df_start, df_stop, on=['eventName', 'x', 'y', 'level'])
# calculate duration
df['duration'] = (df['stop_time'] - df['start_time']).dt.total_seconds()
```

```
Out[87]:
```

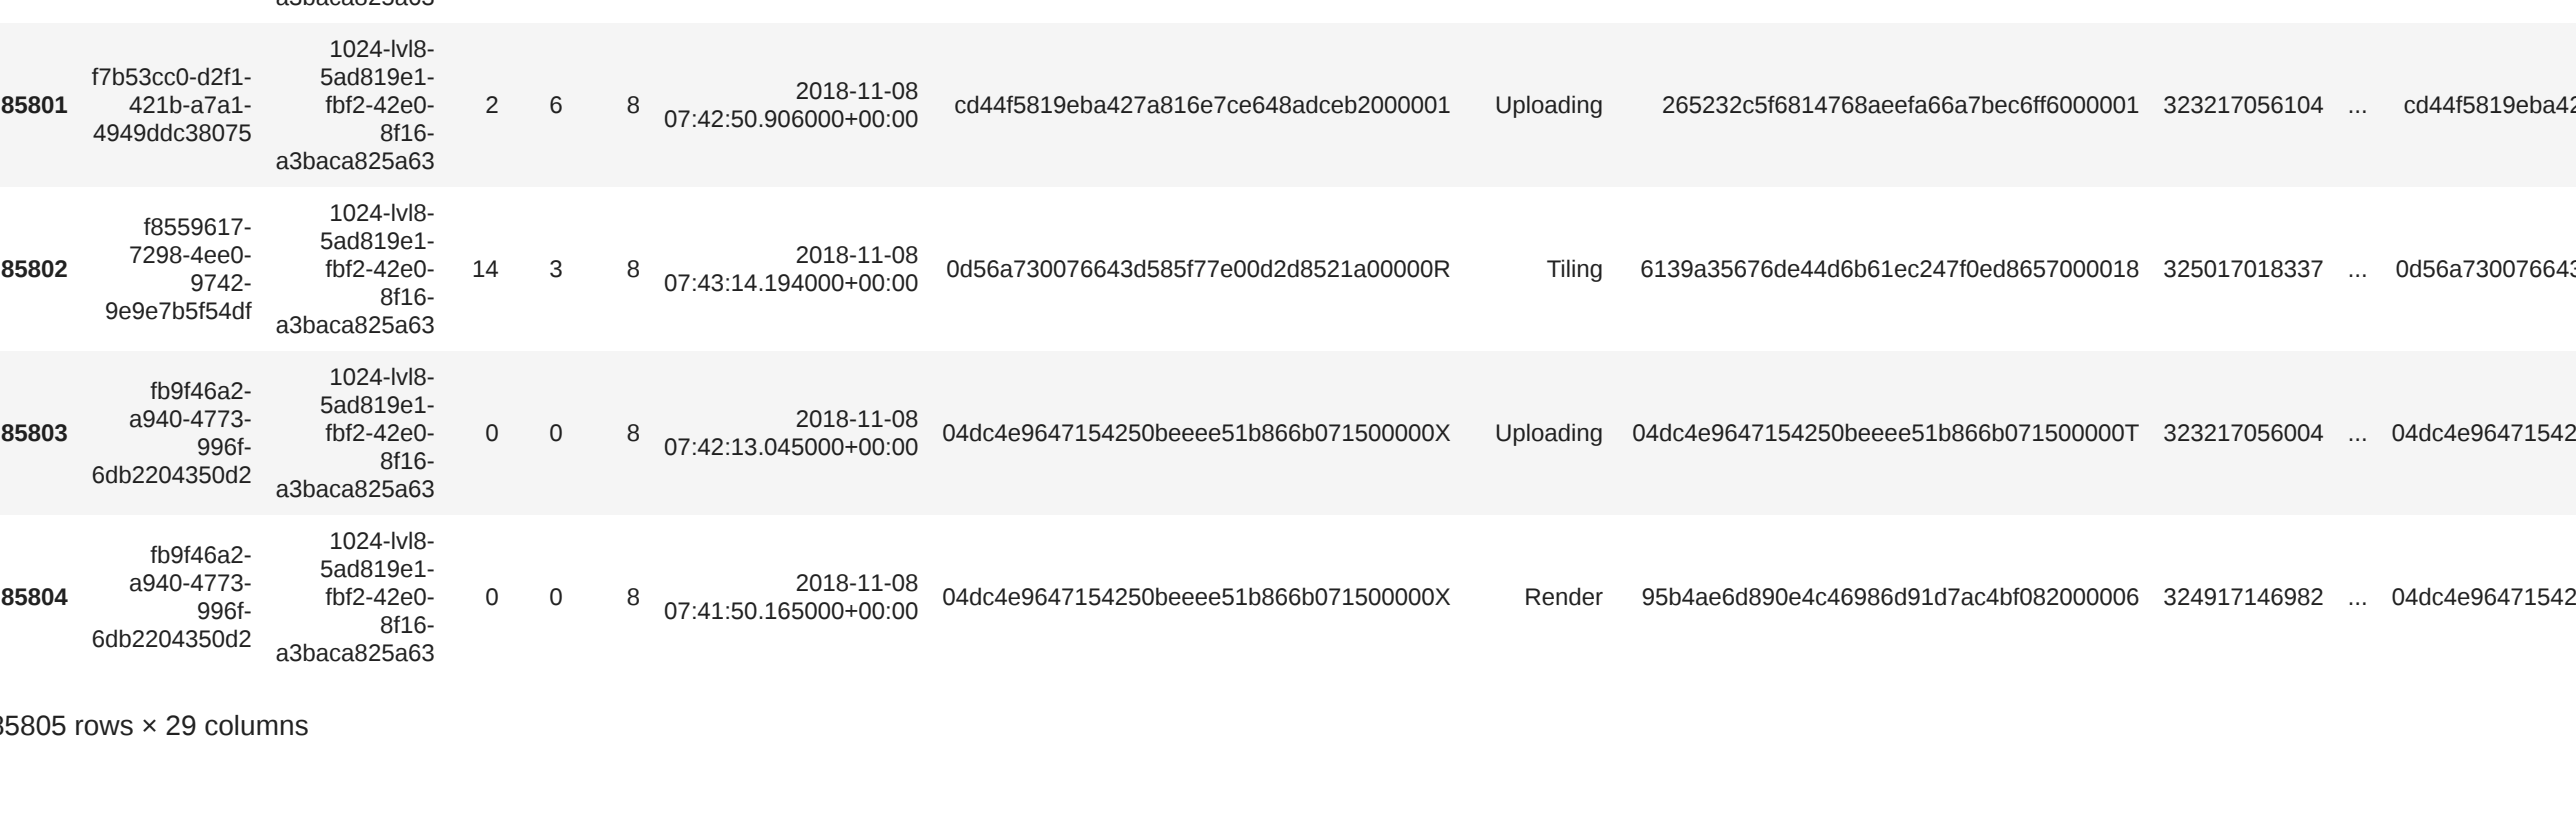
	taskid	jobid	x	y	level	start_time	hostname	x_eventName	eventType	hostname_y	gpuSerial	...
0	00004e77-384c-4802-1346e9f787	1024-M12-7602b3c3-580-48e8-1701-0701	142	190	12	2018-11-08T08:14:48.850000+00:00	83ea61ac1e5427a3bf7b0d041ecaa700001A	Uploading	d8241877c0994572b46e4681e5d144d850000W	323617021323	...	83ea61ac1e5427a3bf7b0d041ecaa700001A
1	00003b80-480a-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	142	86	12	2018-11-08T07:49:17.803000+00:00	83ea61ac1e5427a3bf7b0d041ecaa7000004	Tiling	074591444K046078517041d702622b700001A	325117043985	...	83ea61ac1e5427a3bf7b0d041ecaa7000004
2	00003b80-480a-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	142	86	12	2018-11-08T07:49:17.803000+00:00	83ea61ac1e5427a3bf7b0d041ecaa7000004	Tiling	074591444K046078517041d702622b700001A	325117043985	...	83ea61ac1e5427a3bf7b0d041ecaa7000004
3	00003b80-480a-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	142	86	12	2018-11-08T07:49:17.803000+00:00	83ea61ac1e5427a3bf7b0d041ecaa7000004	Tiling	265232c596814768eaef65a7becf690000Q	323617020801	...	83ea61ac1e5427a3bf7b0d041ecaa7000004
4	00001d4k-1478-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	179	226	12	2018-11-08T08:21:46.0000+00:00	89a1fa7ac27f4e68298607800670700000V	Tiling	89a1fa7ac27f4e68298607800670700000V	324917052134	...	89a1fa7ac27f4e68298607800670700000V
...
85800	ec05a902-7604-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	12	4	8	2018-11-08T07:42:29.161000+00:00	83ea61ac1e5427a3bf7b0d041ecaa7000003	Uploading	8b64beeb87b4c3b20539e81075191800000U	323617021151	...	83ea61ac1e5427a3bf7b0d041ecaa7000003
85801	7f653cd42f1-c21a-17a1-9605-952d364289	1024-M12-7602b3c3-580-48e8-1701-0701	2	6	8	2018-11-08T07:42:29.161000+00:00	c0448519eba427a3bf7b0d041ecaa7000001	Uploading	265232c596814768eaef65a7becf690000Q	323217056104	...	c0448519eba427a3bf7b0d041ecaa7000001
85802	8059617-7708-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	14	3	8	2018-11-08T07:41:14.184000+00:00	0056a730076643b58977f00d0d882c1a0000R	Tiling	6139a3567d6a46861eac217406b8957000018	325017018337	...	0056a730076643b58977f00d0d882c1a0000R
85803	8059617-7708-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	0	0	8	2018-11-08T07:42:13.145000+00:00	046a4c49647154250beeb81b86807150000X	Uploading	046a4c49647154250beeb81b86807150000X	323217060094	...	046a4c49647154250beeb81b86807150000X
85804	8059617-7708-4802-68aa-466a48d85	1024-M12-7602b3c3-580-48e8-1701-0701	0	0	8	2018-11-08T07:41:50.165000+00:00	046a4c49647154250beeb81b86807150000X	Render	95d4a6e8904c468696807f4ac4005200000	324917146882	...	046a4c49647154250beeb81b86807150000X

85805 rows × 29 columns

Which event types dominate task runtimes?

```
In [88]: # Box plot
fig, ax = plt.subplots(figsize=(13, 8))
sns.boxplot(x='eventName', y='duration', data=df, ax=ax)
sns.set(font_scale=1.2)
sns.set(style='darkgrid', palette='gist_earth_r')
ax.set_xlabel('Event Name', fontsize=15)
ax.set_ylabel('Duration of event running (seconds)', fontsize=15)
ax.set_title('Fig 1: Box plot of execution time for each event', fontsize=15)
plt.show()
```

Fig 1: Box plot of execution time for each event



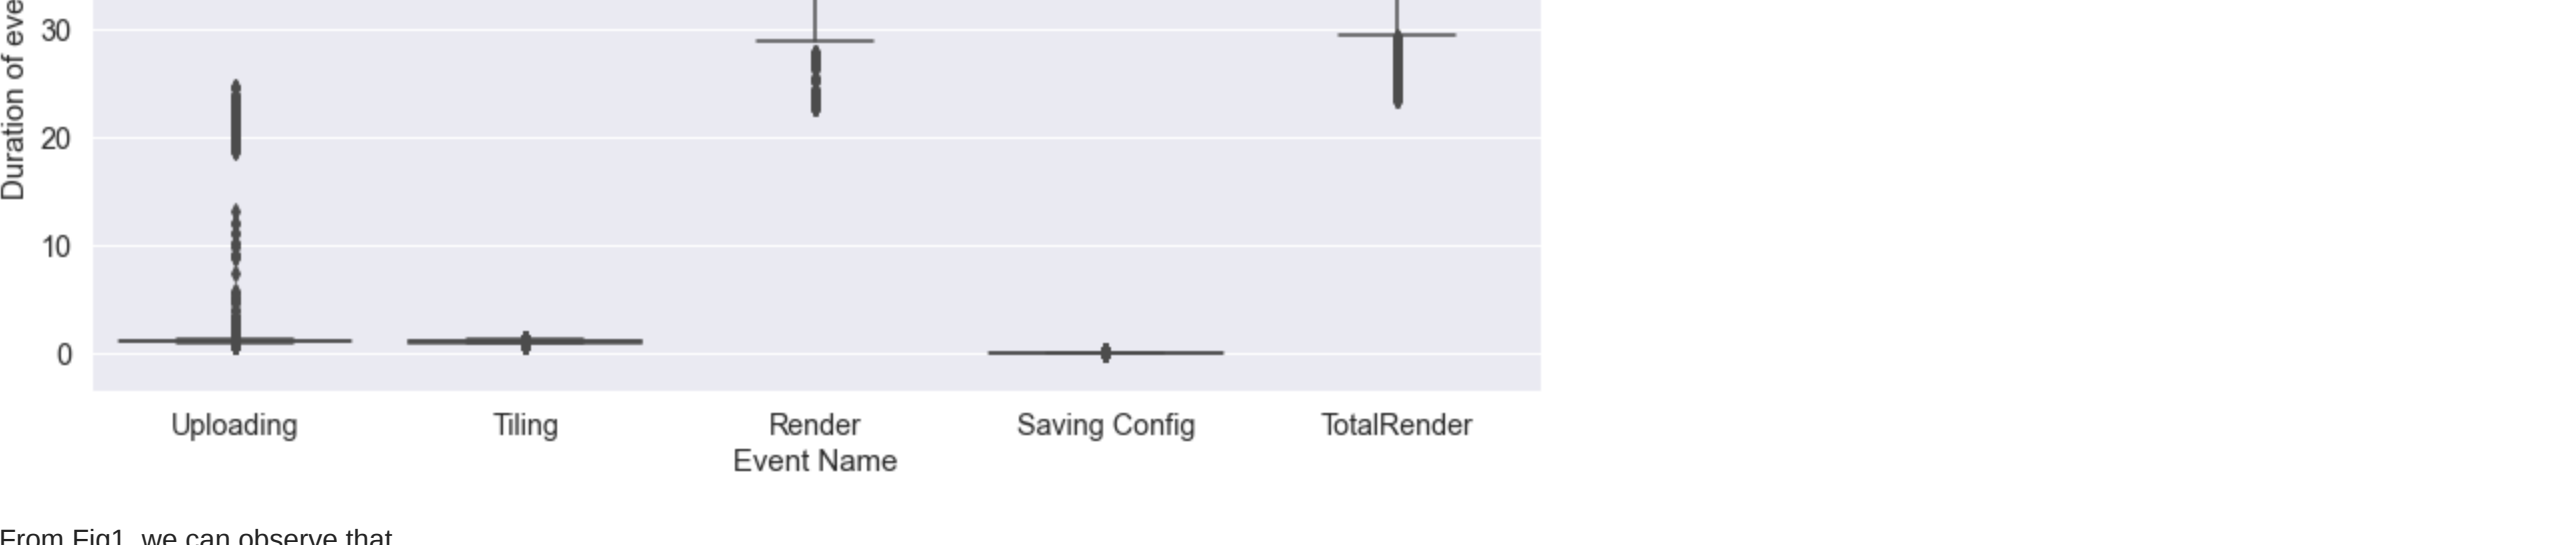
From Fig1, we can observe that

- The majority of the processing time is being consumed by the Event Render function, which is almost equal to the Event TotalRender function. However, there are a significant number of outliers present in both Render and TotalRender.
- On average, rendering takes 42 seconds.
- Saving Config, Tiling, and Uploading do not require a lot of computational time.
- There is also a notable amount of outliers present in the Uploading event.

What is the interplay between GPU temperature and performance?

```
In [89]: # Line plot
fig, ax = plt.subplots(figsize=(12, 7))
ax = sns.lineplot(data=gpudf, x='powerDrawWatt', y='gpuTempC', color='ff7f0e', linewidth=2.5)
ax.set_xlabel('Power Draw by GPU (Watt)', fontsize=15)
ax.set_ylabel('GPU Temperature (Celsius)', fontsize=15)
ax.set_title('Fig 2: Power Draw vs Temperature of the GPU', fontsize=16)
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Fig 2: Power Draw vs Temperature of the GPU



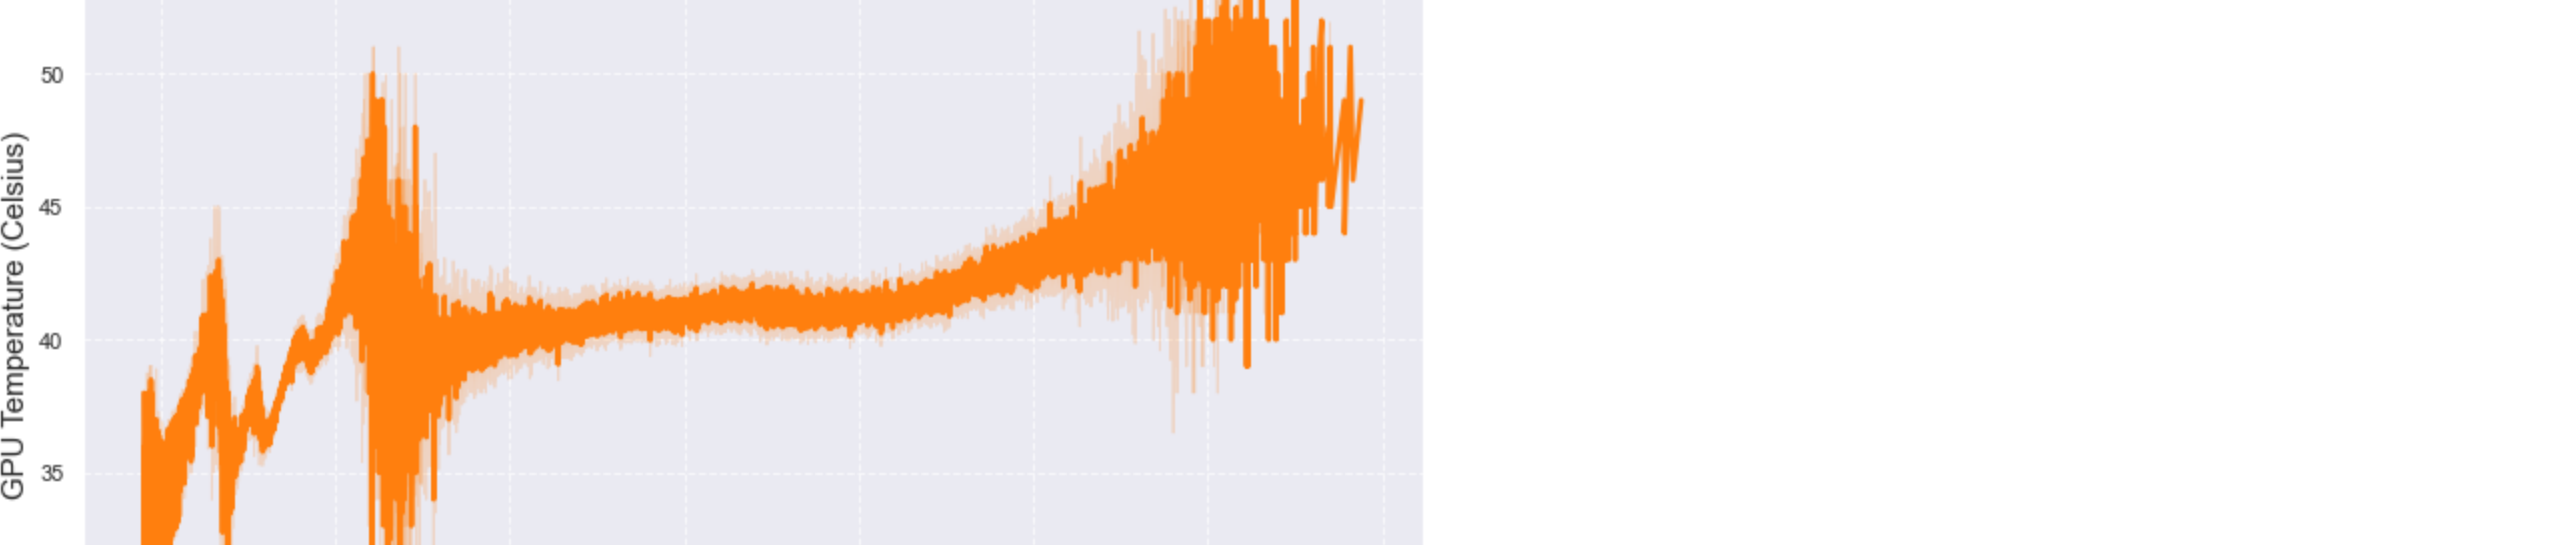
From Fig 2, we can observe that

- As the powerDraw increases the temperature of the gpu also rises.
- We can see that the temperatures fluctuate significantly around 60 watts, indicating that some GPUs with lower performance are reaching higher temperatures at lower power draw.
- We can see that the temperature is increasing with the increase in powerDraw but the increase is not significantly high. This brings us to an conclusion that the cooling infrastructure for the GPU is very good.

As this does not depict the complete picture we will dig more into utilization and memory.

```
In [90]: #line plot
fig, ax = plt.subplots(figsize=(12, 7))
ax = sns.lineplot(data=gpudf, x='powerDrawWatt', y='gpuMemUtilPerc', color='2ca02c', linewidth=2.5)
ax.set_xlabel('Power Draw by GPU (Watt)', fontsize=15)
ax.set_ylabel('GPU Memory Utilility Percentage', fontsize=15)
ax.set_title('Fig 3: Power Draw vs GPU Memory Utilility', fontsize=16)
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Fig 3: Power Draw vs GPU Memory Utilility



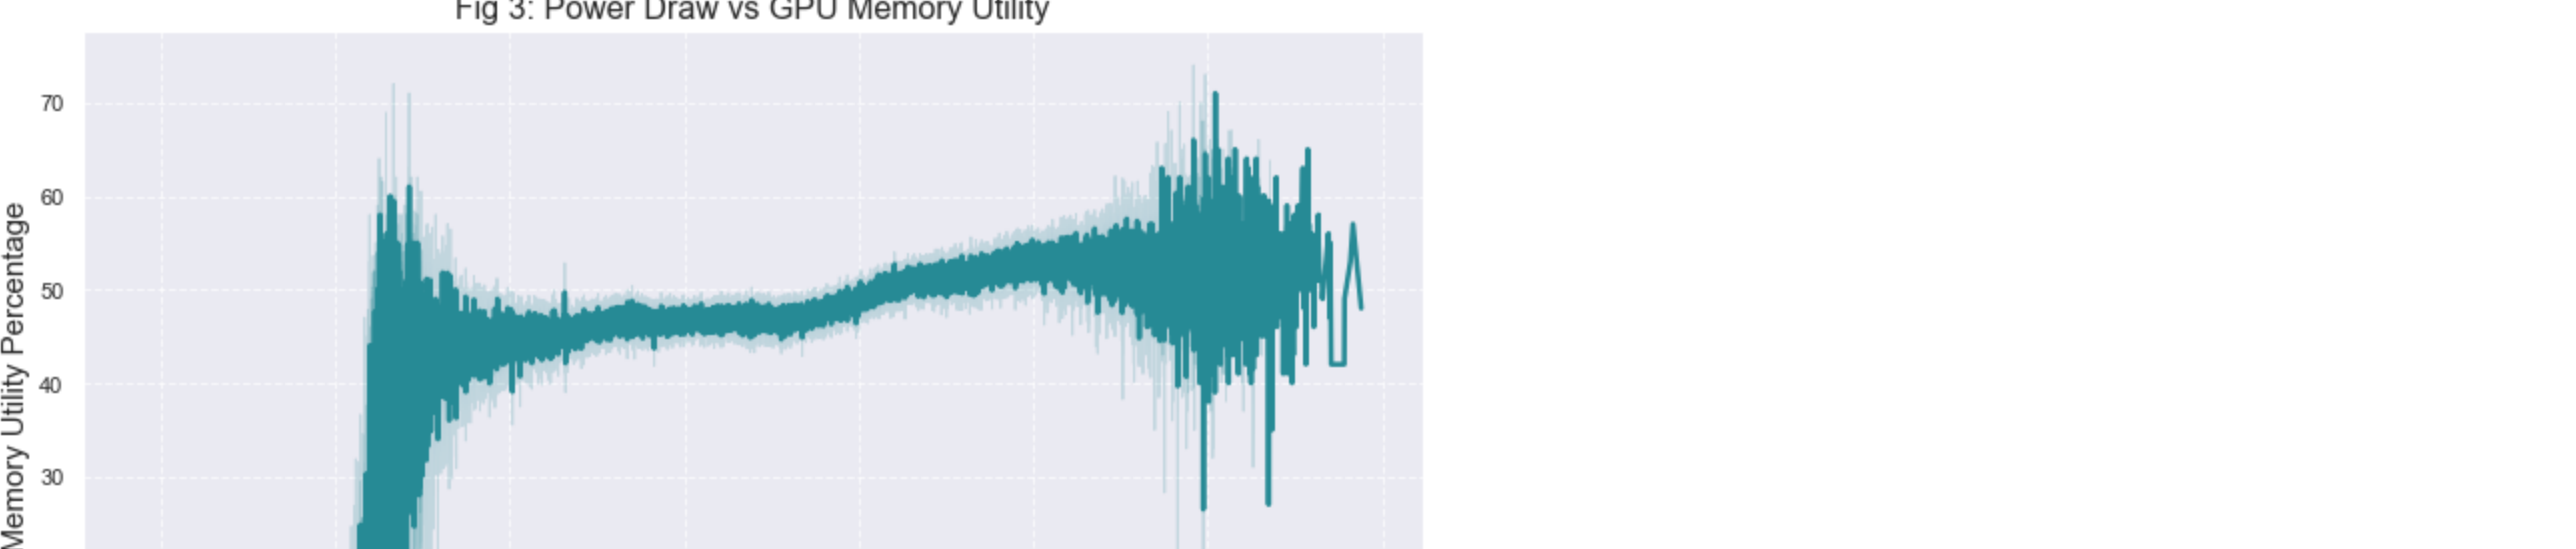
From Fig 3, we can observe that

- Although this graph also depicts the same story, we can see that with the increase in powerDraw more gpu memory is being utilised.
- Same as the previous graph we can see that at 60W few of the GPU's have reached the peak utilization. So we can infer those are low performing GPU's.
- Also at 160W we can see GPU throttling happen, which can result in performance drop.

What is the interplay between increased power draw and render time?

```
In [91]: # Line plot
fig, ax = plt.subplots(figsize=(12, 7))
ax = sns.lineplot(data=gpudf, x='powerDrawWatt', y='duration', markers=True, dashes=False, color='darkblue')
ax.set_xlabel('Power Draw by GPU in Watt', fontsize=15)
ax.set_ylabel('Render Time', fontsize=15)
ax.set_title('Fig 4: Render Time based on Power Draw', fontsize=15)
ax.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Fig 4: Render Time based on Power Draw



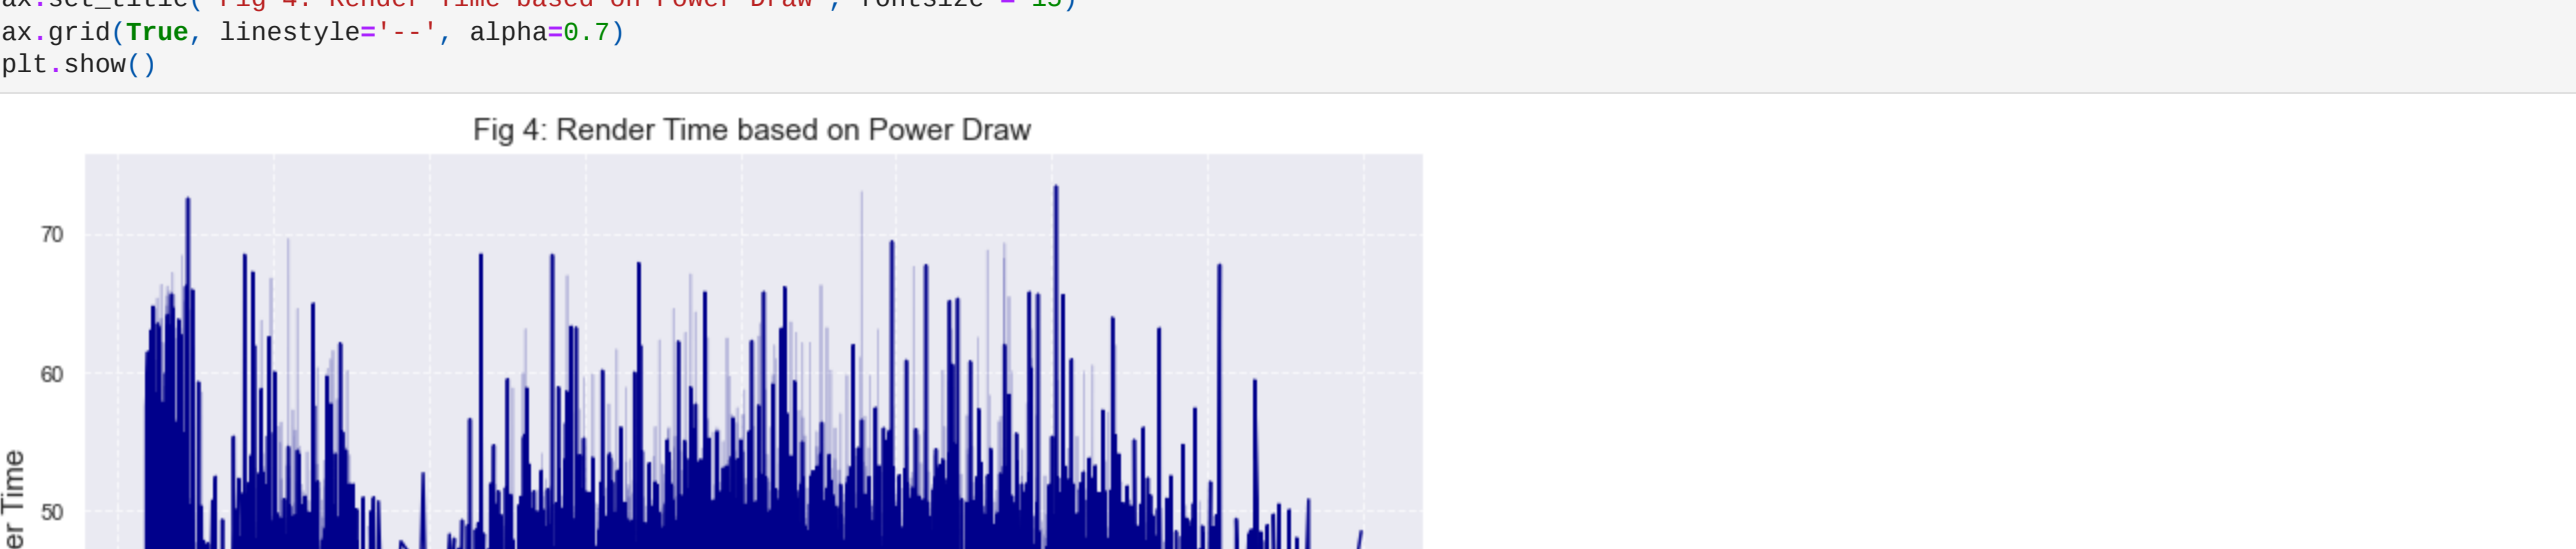
From Fig 4, we can observe that

- We see no increase in render time despite the increase in power draw from the GPUs.
- Although the power demand has changed significantly overall, the render time hasn't improved much, whereas the render time has improved slightly as the temperature has increased.

Can we identify particular GPU cards whose performance differs from other cards?

```
In [92]: colors = "1548fc"
fig, ax = plt.subplots(figsize=(12, 7))
ax = sns.boxplot(x='gpuSerial', y='powerDrawWatt', gpudf[['gpuTempC', 'gpuMemUtilPerc']], patch_artist=True, labels = gpu_stats, boxprops = dict(facecolor='lightcoral', edgecolor='black', linewidth=2.5))
ax.set_xlabel('GPU Statistics', fontsize=15)
ax.set_ylabel('Values', fontsize=15)
ax.grid(True, linestyle='--', alpha=0.7)
plt.rcParams['figure.figsize'] = (12, 7)
plt.show()
```

Fig 5: Box Plot of GPU Statistics



From Fig 5, we can observe that

- The average power consumption is around 95 watts.
- The average temperature is 35 degrees Celsius.
- The average GPU utilization is 80%.
- The average memory utilization is 40%.

This implies that there is a roughly equal number of underperforming and high-performing GPUs, as they all seem to be maintaining an average performance level at high GPU utilization.

```
In [93]: gpudf['gpuSerial'] = gpudf['gpuSerial'].astype(str)
gpu_powerdraw = gpudf.groupby('gpuSerial', as_index=False)['powerDrawWatt'].mean()
```

```
In [94]: # Line plot
fig, ax = plt.subplots(figsize=(12, 7))
sns.set(style='whitegrid')
plot = sns.lineplot(data=gpu_powerdraw,
```